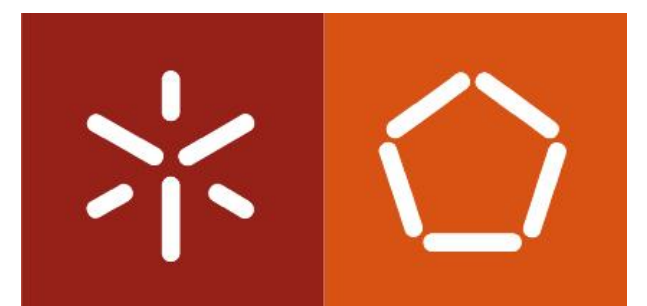


Operating Systems

(Sistemas Operativos)

Guide 1: Files




File Interface

The File Descriptor (FD)

- **File Descriptor (FD)**
 - integer that represents files
 - used for reading/writing files, pipes, or I/O devices
- **Array of file descriptors** (or file descriptor table)
 - Maintained by each process
 - Starts with three **open standard FDs**
 - 0 - Standard input:** read from keyboard
 - 1 - Standard output:** write to screen
 - 2 - Standard error:** write error messages

FD array of Process A



File Interface

Reading and Writing

#include <unistd.h>

- `ssize_t read(int fd, void *buf, size_t nbyte)`
 - **fd**: the file descriptor
 - **buf**: buffer from where the content is read
 - **nbyte**: max number of bytes to read¹
 - Returns: **number of bytes read** (or -1 on errors)

For more information: `$ man 2 read`

#include <unistd.h>

- `ssize_t write(int fd, const void *buf, size_t nbyte)`
 - **fd**: the file descriptor
 - **buf**: buffer with content to be written
 - **nbyte**: number of bytes to write from the buffer¹
 - Returns: **number of bytes written** (or -1 on errors)

For more information: `$ man 2 write`

¹**Buffer overrun**: number of bytes to read/write may not be larger than the memory allocated for the buffer!

File Interface

Opening files

#include <fcntl.h>

- *int* **open**(*const char *pathname, int oflag [, mode]*)
 - **pathname**: absolute or relative pathname
 - **oflag** - opening mode(s)
 - *O_WRONLY*, *O_RDONLY*, *O_RDWR* - open for writing, reading, or both, respectively
 - *O_CREAT* - create file if it does not exist
 - *O_TRUNC* - truncate file to zero length
 - *O_APPEND* - write data to the end of the file
 - **mode** - file permissions (required with *O_CREAT*)
 - *0600* - owner of the file can read/write
 - Returns: **file descriptor** (or -1 on errors)

For more information: *\$ man 2 open*

Note: Do not forget to close file descriptors when these are no longer needed (see Slide 15)

File Interface

Example: opening a file

```
int fd = open("foo.txt", O_CREAT | O_TRUNC | O_RDWR, 0600)  
fd = 3
```

FD array of process A

0	→	stdin
1	→	stdout
2	→	stderr
3	→	

Entry at the open file table

MODE	RW
OFFSET	0
#REF	1
INODE	


inode for file foo.txt

I-NUMBER	100
SIZE	0
#REF	1
...	

File Interface

Example: writing a file

```
ssize_t bytes_written = write(fd, "abcde", 5)  
bytes_written = 5
```



FD array of process A

0	→	stdin
1	→	stdout
2	→	stderr
3	→	

Entry at the open file table

MODE	RW
OFFSET	5
#REF	1
INODE	

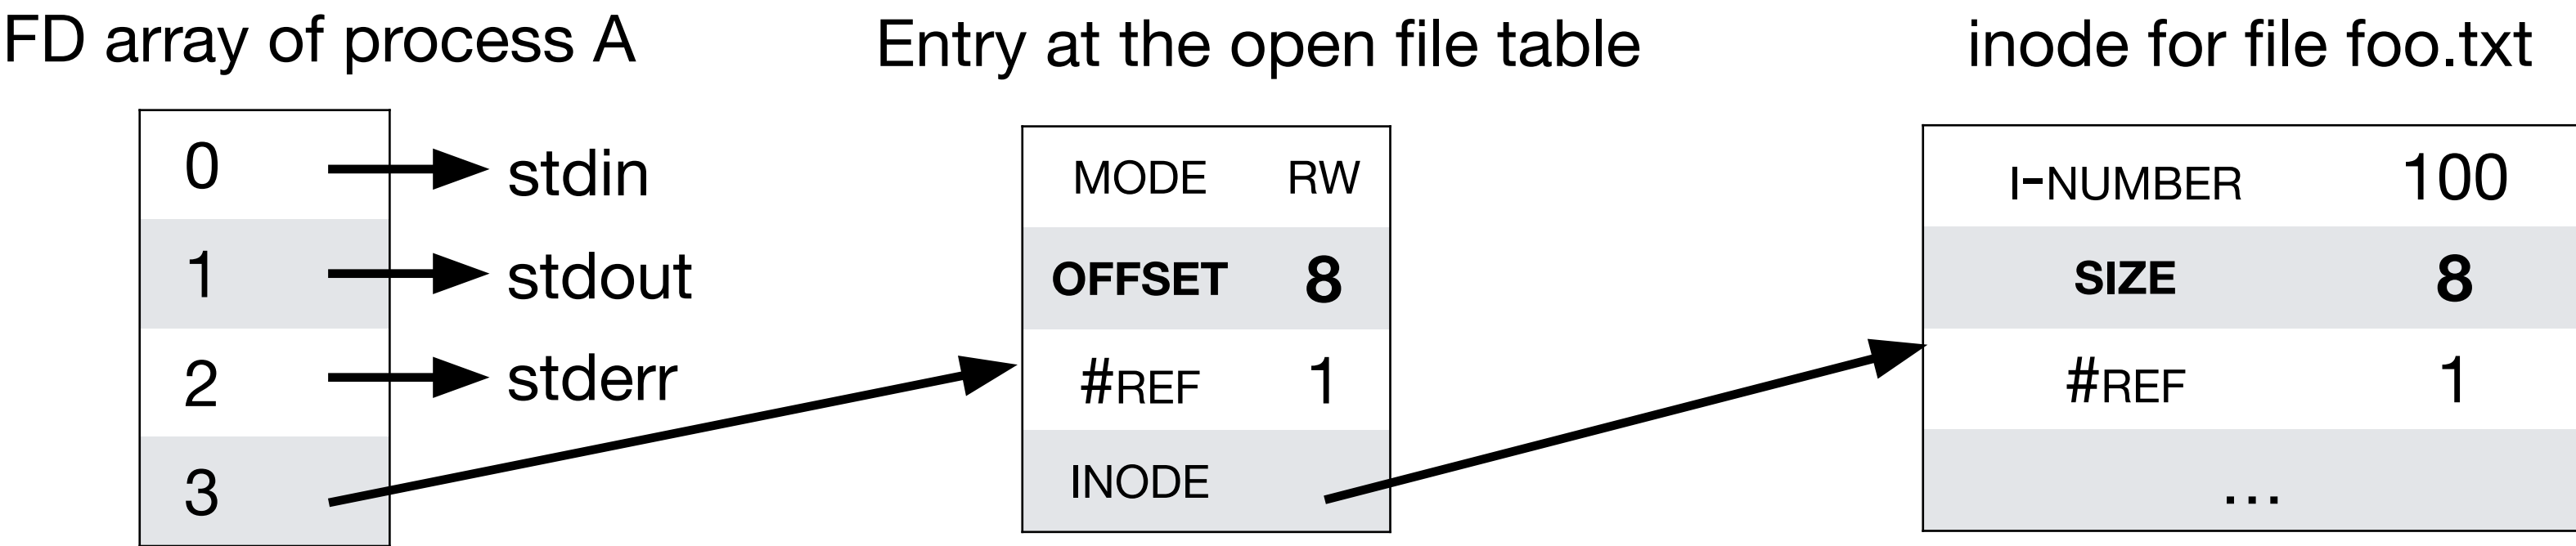

inode for file foo.txt

I-NUMBER	100
SIZE	5
#REF	1
...	

File Interface

Example: writing a file

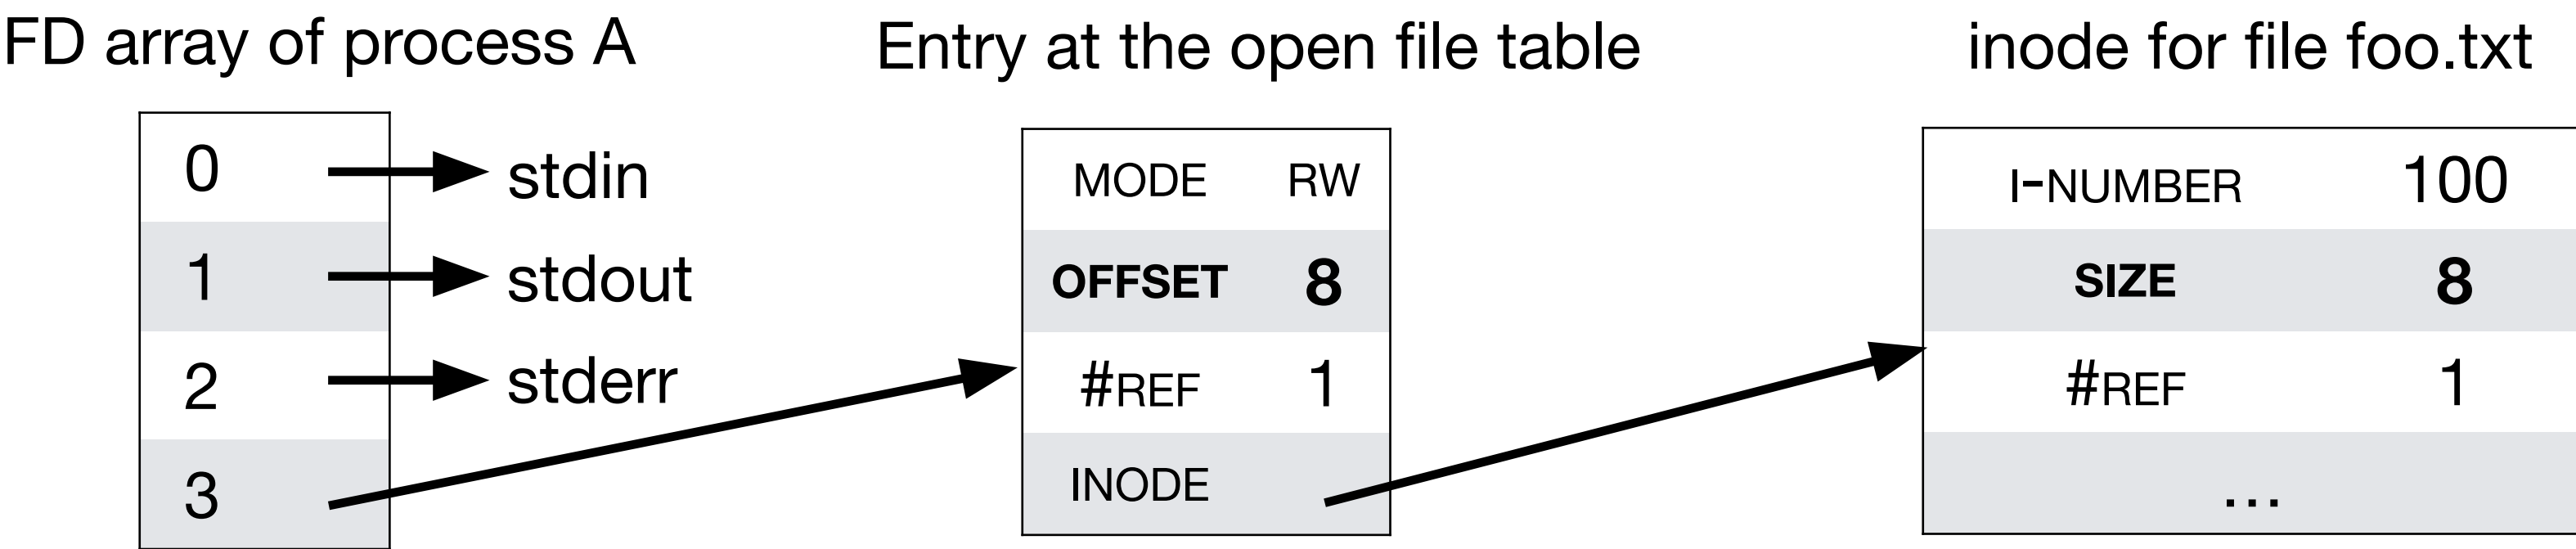

```
bytes_written = write(fd, "fgh", 3)  
bytes_written = 3
```



File Interface

Example: reading a file


```
bytes_read = read(fd, buf, 8)
bytes_read = ?
```



File Interface

Example: reading a file

```
bytes_read = read(fd, buf, 8)
bytes_read = 0
```



FD array of process A

0	→	stdin
1	→	stdout
2	→	stderr
3	→	

Entry at the open file table

MODE	RW
OFFSET	8
#REF	1
INODE	

inode for file foo.txt

I-NUMBER	100
SIZE	8
#REF	1
...	

It will **return 0**, the program is reading (the offset is positioned) at the **end of the file!**

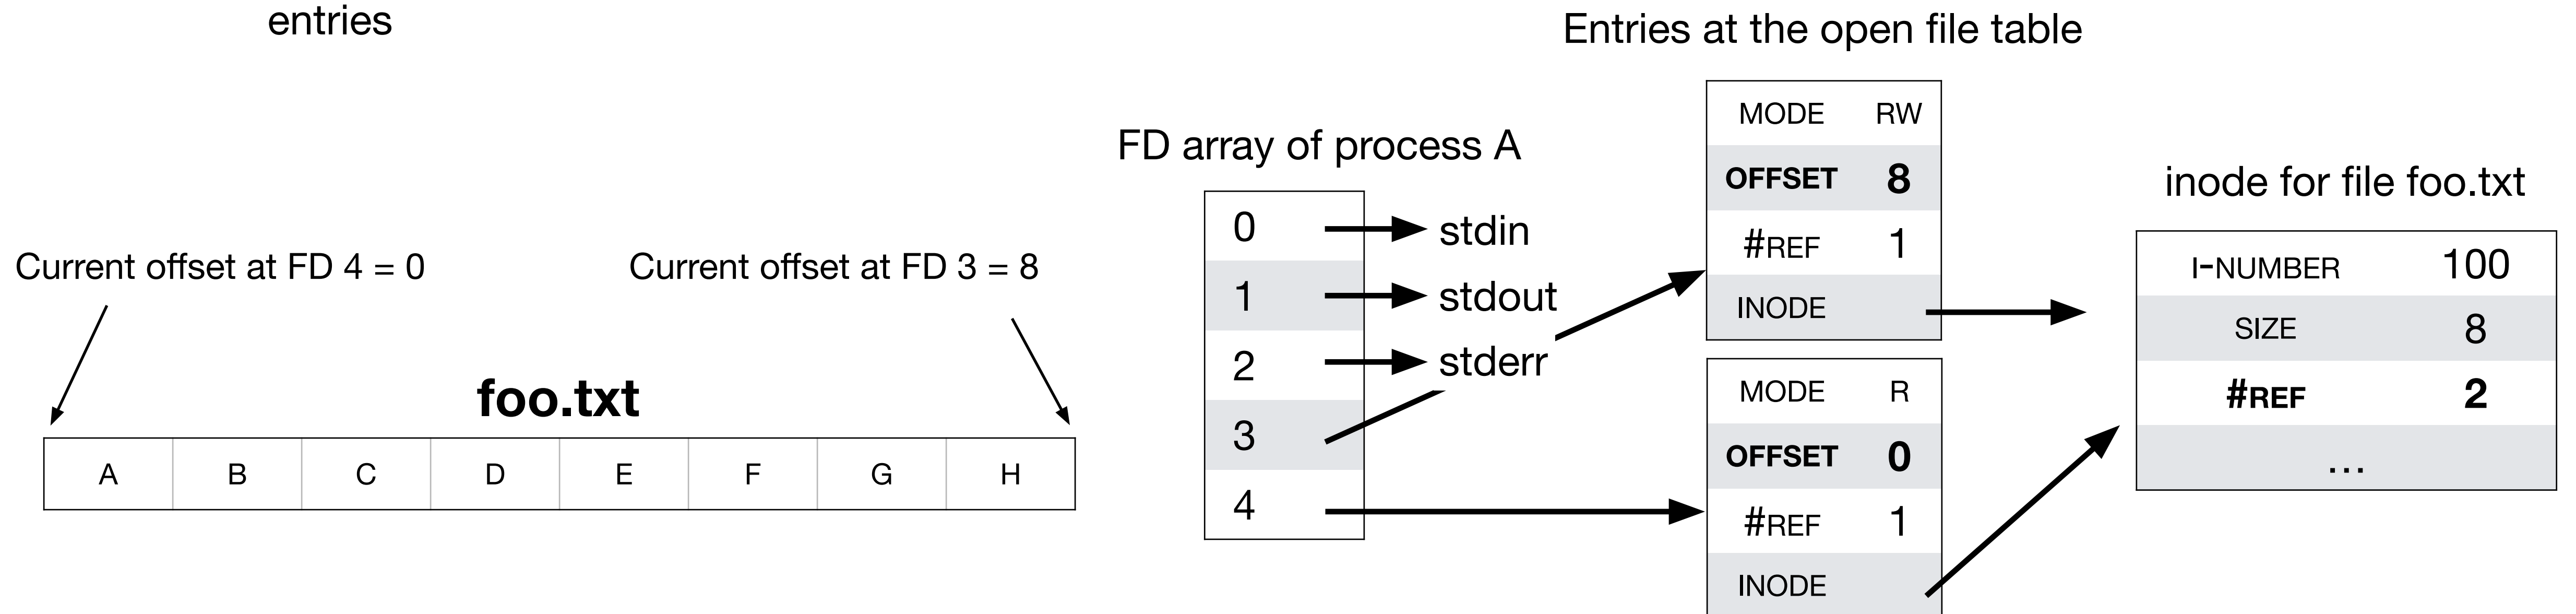
File Interface

Multiple file descriptors

- **Option 1:** Using another descriptor

```
int fd1 = open("foo.txt", O_RDONLY)
```

- new entry at the **open file table**, with offset = 0
- `read(fd1, ...)`
- Note: the **inode** is **shared** across the two file table entries



File Interface

Random file access

- **Option 2:** Do a random access to the descriptor's offset 0

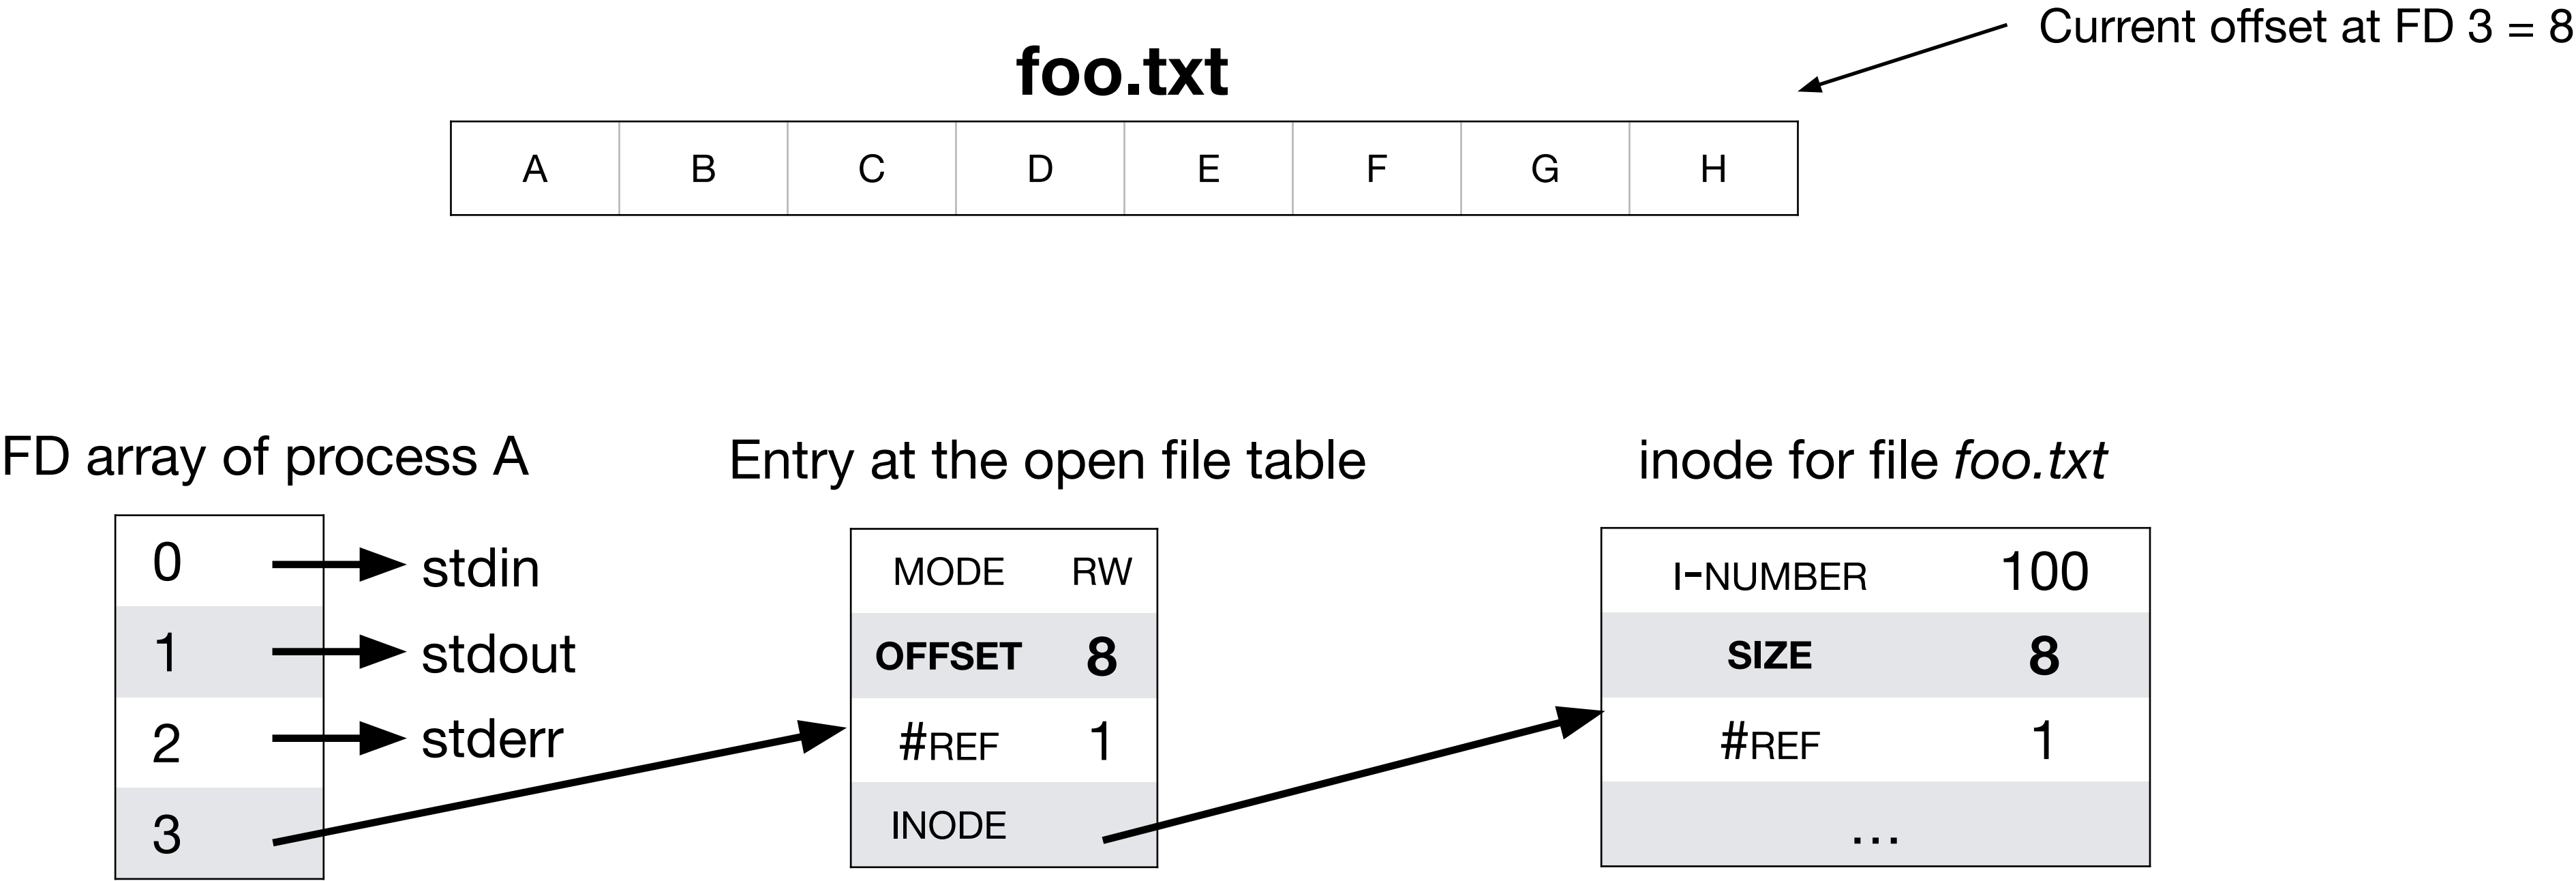
#include <unistd.h>

- *off_t* ***lseek***(*int fd*, *off_t offset*, *int whence*)
 - **fd:** the file descriptor
 - **offset:** the number of bytes to move forward or backward (it can be negative)
 - **whence:** from where to move:
 - SEEK_SET - from the beginning of the file
 - SEEK_END - from the end of the file
 - SEEK_CUR - from the current offset
 - Returns: the **resulting offset at the file** or -1 on **error**

For more information: *\$ man 2 lseek*

File Interface

Example: seeking and reading a file



File Interface

Example: seeking and reading a file

```
off_t res_off = lseek(fd, SEEK_SET, 0)  
res_off = 0
```

Current offset at FD 3 = 0



FD array of process A

0	→	stdin
1	→	stdout
2	→	stderr
3	→	

Entry at the open file table

MODE	RW
OFFSET	0
#REF	1
INODE	

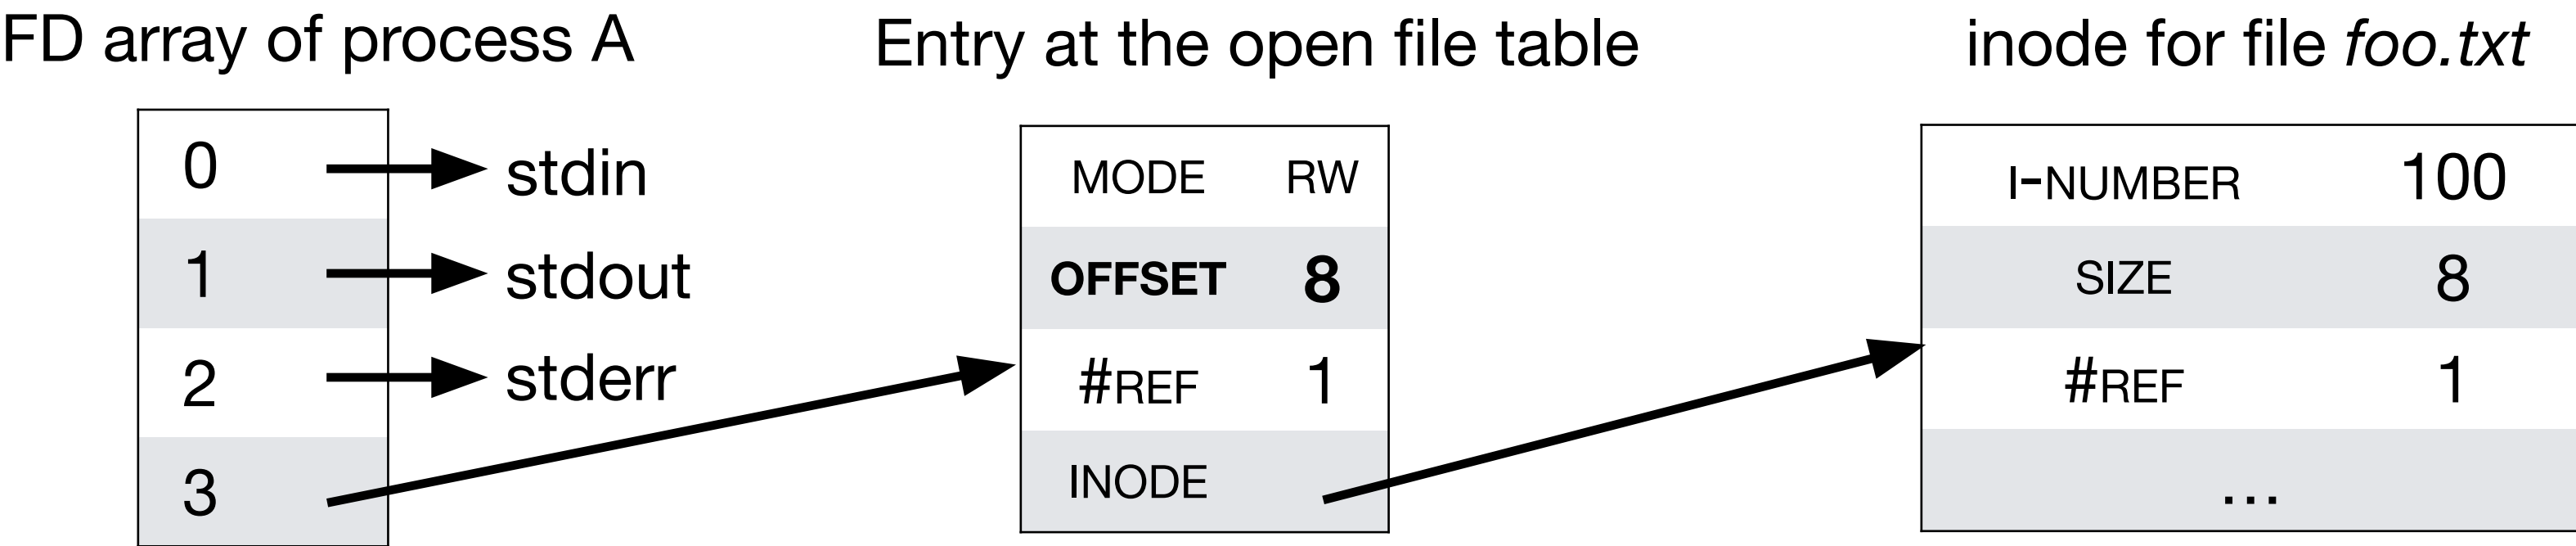

inode for file *foo.txt*

I-NUMBER	100
SIZE	8
#REF	1
...	

File Interface

Example: seeking and reading a file

```
bytes_read = read(fd, buf, 8)
bytes_read = 8
```



File Interface

Example: closing and deleting files

#include <unistd.h>

- **int close(int fd)**
 - **fd:** the file descriptor
 - Returns: 0 on success or -1 on error

FD array of process A

0	→	stdin
1	→	stdout
2	→	stderr
3	→	✗

inode for file *foo.txt*

I-NUMBER	100
SIZE	8
#REF	0
...	

inode is not deleted!

For more information: `$ man 2 close`

#include <unistd.h>

- **int unlink(const char* pathname)**
 - **pathname:** absolute or relative pathname
 - Returns: 0 on success or -1 on error

deletes the inode when its *#ref* reaches 0

- Remember that multiple unrelated processes may have the same file opened

For more information: `$ man 2 unlink`